Gradient Boosting

Elena Kantonistova

HSE, 2021

Gradient Boosting & Practice

Theoretical part:

- Classical Boosting

- Popular gradient boosting implementations

Practical part - Python notebook with:

- Feature engineering

- Feature selection

- Modeling

Gradient boosting

- Gradient boosting is a machine learning algorithm for classification and regression.

- It is an ensemble of multiple weak learners b_i:

$$a(x) = \sum_{i=1}^{N} b_i(x)$$

- In gradient boosting, we build trees step by step and add them to the composition.

Gradient boosting

Key idea: at each next step, we look for a basic algorithm that corrects the composition error in the previous step
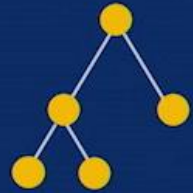
Boosting algorithm

x is an object and y is a target vector.

- Boosting builds an ensemble of trees *one-by-one*, then the predictions of the individual trees *are summed*.

- If an ensemble has two trees at the current step, we have:
$$a(x) = b_1(x) + b_2(x)$$

- *The next decision tree $b_3(x)$ tries to predict the error of the composition.* That means that the target vector for the next algorithm is

$$y_{new} = y - a(x)$$

Boosting algorithm

1.E

dt1

error: y - dt1(x)

2.E

*dt2*

add dt2 to composition

Boosting: a(x) = dt1(x) + dt2(x)

2.D

Gradient boosting overfits

- Since at each iteration we reduce the error on the training data, *at some point boosting will begin to overfit* - the error on the test data will begin to grow.

- *You need to find optimal number of trees* to minimize test error.
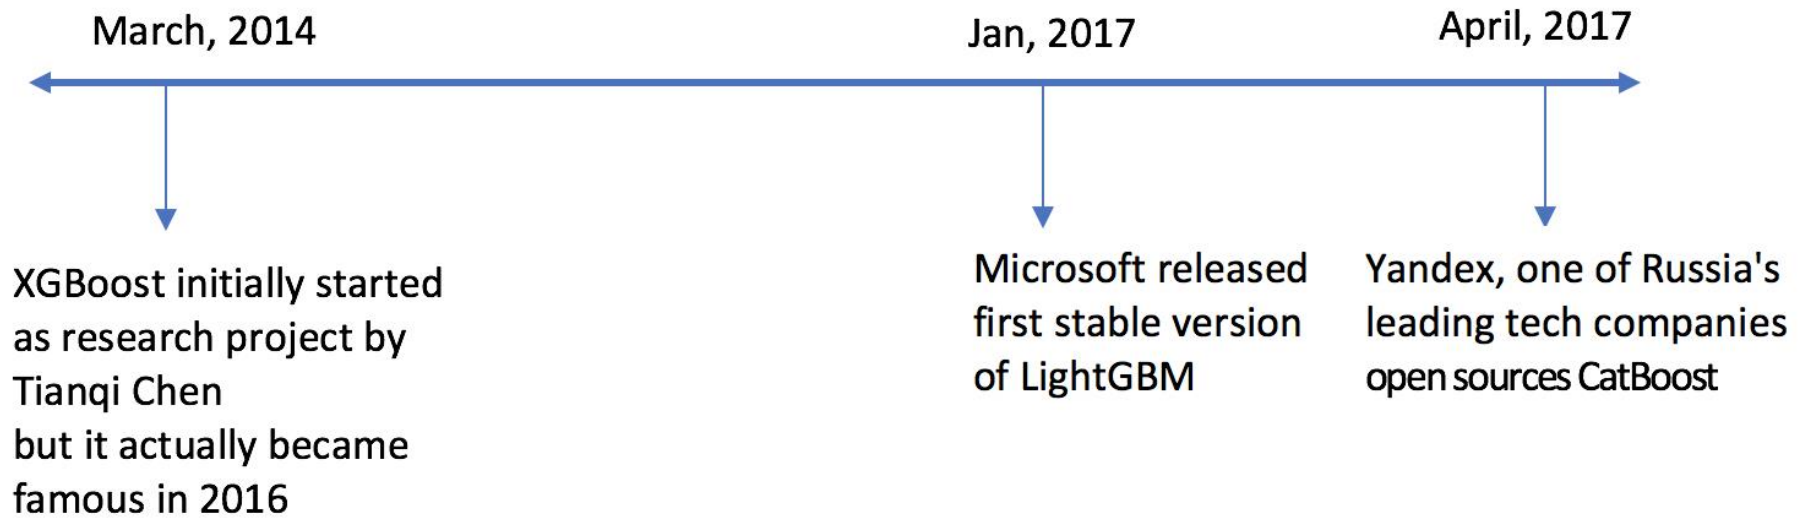
Gradient Boosting implementations

XGBoost

CatBoost

LightGBM

XGBoost, LightGBM, CatBoost

March, 2014

Jan, 2017

April, 2017

XGBoost initially started
as research project by
Tianqi Chen
but it actually became
famous in 2016

Microsoft released
first stable version
of LightGBM

Yandex, one of Russia's
leading tech companies
open sources CatBoost

XGBoost loss (Extreme gradient boosting)

Main differences between XGBoost and classic gradient boosting:

1. when constructing each weak learner, the *loss function is approximated to second-order* (in gradient boosting – to first order)

XGBoost loss (Extreme gradient boosting)

Main differences between XGBoost and classic gradient boosting:

2. we add a *regularization* to the loss: we penalize the tree

- for the number of leaves

- for large forecasts in the leaves

XGBoost

- [Xgboost documentation](#)
- https://machinelearningmastery.com/gentle-introduction-xgboost-applied-machine-learning/
- https://www.datacamp.com/community/tutorials/xgboost-in-python

CATBOOST

CatBoost is the first Russian machine learning algorithm developed to be open source. The algorithm was developed in the year 2017 by machine learning researchers and engineers at Yandex (a technology company).

CATBOOST

Main advantages:

- Very good quality on many datasets
- High prediction speed
- Support for both numerical and categorical features
- Fast GPU support

CATBOOST

- CatBoost builds *symmetric trees*:

Symmetric trees

Height > 170 — Weight > 65 — Age > 16

Height > 170 — Weight > 65, Weight > 65

CATBOOST

Symmetric trees (ST)

Height > 170

Weight > 65

Age > 16

Height > 170

Weight > 65   Weight > 65

- ST overfit less then standard trees because of the fixed structure
- Thanks to ST, CatBoost shows good quality even with default parameters

CATBOOST

- CatBoost *encodes categorical features*.

Categorical features

Categorical data

Occupation → Engineer, Designer, Writer, Manager, HR

CATBOOST

- CatBoost encodes categorical features.

Categorical features support

> One-hot encoding

> Statistics based on category and category plus label value

> Usage of several permutations

> Greedy constructed feature combinations

$$i \longrightarrow \frac{1 + 1 + 0 + a * Prior}{3 + a}$$

CATBOOST

- CatBoost encodes categorical features.

Categorical features support

> One-hot encoding

> Statistics based on category and category plus label value

> Usage of several permutations

> Greedy constructed feature combinations

|	SDE		1
	SDE		1
	SDE		0
	PR		
i	SDE		1
	PR		

$$i \longrightarrow \frac{1 + 1 + 0 + a * Prior}{3 + a}$$

- You don't need to spend time to find optimal set of encodings for your dataset. CatBoost does all the job.

CATBOOST

- CatBoost *uses ordered boosting*

Ordered boosting

$$\text{leafValue(doc)} = \sum_{i=1}^{doc} \frac{g(\text{approx}(i), \text{target}(i))}{docs\ in\ the\ past}$$

- DT overfits because we use the same data to find tree structure and to make the predictions in leaves

- Ordered boosting is a procedure that reduces overfitting by adding randomness in the process of making the predictions in leaves

CatBoost

- [CatBoost documentation](#)
- [video presentation (very good!)](#)
- https://coderzcolumn.com/tutorials/machine-learning/catboost-an-in-depth-guide-python

LightGBM

LightGBM was developed by researchers in Microsoft Research in 2017. It is also an open source library.

LightGBM

Main advantages of LightGBM:

- "Light" = high speed

- can handle large size of data

- focuses on the accuracy of results

- supports GPU learning

LightGBM

Usually boosting grows tree *level-wise*:

Level-wise tree growth

LightGBM

LightGBM grows tree *leaf-wise*:

Leaf-wise tree growth

- it will choose the leaf with max delta loss to grow
- leaf-wise algorithm can reduce more loss than a level-wise algorithm

LightGBM

LightGBM works fast because of *feature bucketing*.

• it uses a histogram based algorithm
to find the optimal split point
while creating a weak learner

• therefore, each continuous numeric
feature should be split into
discrete bins

An example of how binning can reduce the number of splits to explore. The features must be sorted in advance for this method to be effective.

LightGBM

- LightGBM documentation
- https://towardsdatascience.com/what-makes-lightgbm-lightning-fast-a27cf0d9785e
- https://medium.com/@pushkarmandot/https-medium-com-pushkarmandot-what-is-lightgbm-how-to-implement-it-how-to-fine-tune-the-parameters-60347819b7fc

Thank you for your attention!

Now it's time for practice!